


DAY – 9**TOPIC: PostgreSQL Upgrade & Uninstallation Process**

Detailed guide to upgrading PostgreSQL, covering why upgrades matter, how to prepare, and step-by-step instructions for both minor and major version upgrades.

🚀 Why Upgrade PostgreSQL?

Upgrading your PostgreSQL database is essential for:

- **Security:** New versions patch vulnerabilities.
- **Performance:** Improved indexing, query planning, and memory usage.
- **Features:** New data types (e.g., jsonb), SQL enhancements, and better concurrency.
- **Support:** Older versions reach End of Life (EOL), losing updates and fixes.

⌚ Types of PostgreSQL Upgrades



Upgrade Type	Description	Method
Minor	Patch-level updates (e.g., 16.4 → 16.5)	Replace binaries, restart server
Major	Version jump (e.g., 14 → 16)	Use pg_upgrade, pg_dump, or replication

Minor upgrades are simple and backward-compatible. Major upgrades may involve changes to internal storage formats and system catalogs.

Pre-Upgrade Checklist

- 1. Read Release Notes** Review all changes and migration notes for the target version.
- 2. Check Extension Compatibility** Ensure all extensions (e.g., PostGIS, pg_stat_statements) support the new version.
- 3. Test in a Staging Environment** Validate application behavior and performance before production rollout.
- 4. Backup Your Data** Use:

`pg_dumpall > alldb.bak`

`pg_dump dbname > dbname.bak`

- 5. Plan for Downtime** Schedule during off-peak hours or use replication for near-zero downtime.

Step-by-Step: Major Upgrade Using pg_upgrade

1. Install New PostgreSQL Version

Download and install the new binaries from postgresql.org.

2. Initialize New Cluster

`/usr/localpgsql/bin/initdb -D /usr/localpgsql/data_new`

3. Run pg_upgrade

```
/usr/localpgsql/bin/pg_upgrade \
-b /usr/lib/postgresql/14/bin \
-B /usr/lib/postgresql/16/bin \
```



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

```
-d /var/lib/postgresql/14/main \
-D /var/lib/postgresql/16/main \
-o "-c config_file=/etc/postgresql/14/main/postgresql.conf" \
-O "-c config_file=/etc/postgresql/16/main/postgresql.conf"
```

4. Analyze and Reindex

```
./analyze_new_cluster.sh
```

```
./delete_old_cluster.sh
```

💡 Alternative: Dump and Restore

Use when pg_upgrade isn't viable (e.g., cross-platform migration):

```
pg_dump -Fc dbname > dbname.dump
```

```
pg_restore -d newdb dbname.dump
```

⌚ Near-Zero Downtime Option

Use logical replication:

1. Set up a replica on the new version.
2. Stream changes using pglogical or native replication.
3. Switch traffic once synced.

Learn more

✖️ Post-Upgrade Tasks

- Verify application functionality.
- Monitor logs for errors or performance issues.
- Keep the old version temporarily for rollback if needed.

⚠️ Common Challenges in PostgreSQL Upgrades

1. Deprecated Features

- **Issue:** Older PostgreSQL versions may rely on features that are removed or changed in newer releases.
- **Example:** Implicit casts removed in PostgreSQL 8.3 broke many legacy queries.



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

- **Solution:** Review release notes and test queries for compatibility before upgrading.

2. Downtime Risk

- **Issue:** Upgrades can cause service interruptions.
- **Solution:** Use pg_upgrade or logical replication for minimal downtime. Always schedule upgrades during off-peak hours.

```
pg_upgrade -b /usr/local/pgsql.old/bin -B /usr/local/pgsql.new/bin \
-d /usr/local/pgsql.old/data -D /usr/local/pgsql.new/data \
-p 5432 -P 5433
```

3. Compatibility Issues

- **Issue:** SQL behavior, data types, or extensions may change.
- **Example:** JSON vs. JSONB—older versions stored JSON as plain text, while newer versions validate structure and support indexing.
- **Solution:** Validate application logic and extension compatibility in a staging environment.

4. Data Loss

- **Issue:** Improper upgrade procedures can corrupt or delete data.
- **Solution:** Always back up using pg_dump, pg_dumpall, or filesystem snapshots. Test restore procedures.

5. Performance Regressions

- **Issue:** New versions may introduce changes that degrade performance for certain workloads.
- **Solution:** Benchmark queries before and after upgrade. Tune configuration parameters for the new version.

Benefits of Upgrading PostgreSQL

Upgrading unlocks powerful features and performance improvements. For example, PostgreSQL 13 introduced:

- **Improved Query Planning:** Better statistics for complex filters.
- **Parallel Vacuuming:** Faster table maintenance.

VACUUM (PARALLEL 4) customers;



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

- **Incremental Sorting:** Efficient multi-key sorting.

`SELECT name FROM employees ORDER BY salary DESC, name ASC;`

Step-by-Step Upgrade Guide

Here's a high-level upgrade workflow:

◆ Preparation

1. Read release notes
2. Audit extensions and custom functions
3. Back up your data:

`pg_dumpall > alldb.bak`

◆ Installation

4. Install new PostgreSQL binaries
5. Initialize new cluster:

`initdb -D /new/data`

◆ Migration

6. Use pg_upgrade:

`pg_upgrade -b /old/bin -B /new/bin -d /old/data -D /new/data`

7. Run post-upgrade scripts:

`./analyze_new_cluster.sh`

`./delete_old_cluster.sh`

◆ Validation

8. Test application behavior
9. Monitor logs and performance
10. Keep old version temporarily for rollback

Common Challenges in PostgreSQL Upgrades

1. Deprecated Features



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

- **Issue:** Older PostgreSQL versions may rely on features that are removed or changed in newer releases.
- **Example:** Implicit casts removed in PostgreSQL 8.3 broke many legacy queries.
- **Solution:** Review release notes and test queries for compatibility before upgrading.

2. Downtime Risk

- **Issue:** Upgrades can cause service interruptions.
- **Solution:** Use pg_upgrade or logical replication for minimal downtime. Always schedule upgrades during off-peak hours.

```
pg_upgrade -b /usr/local/pgsql.old/bin -B /usr/local/pgsql.new/bin \
-d /usr/local/pgsql.old/data -D /usr/local/pgsql.new/data \
-p 5432 -P 5433
```

3. Compatibility Issues

- **Issue:** SQL behavior, data types, or extensions may change.
- **Example:** JSON vs. JSONB—older versions stored JSON as plain text, while newer versions validate structure and support indexing.
- **Solution:** Validate application logic and extension compatibility in a staging environment.

4. Data Loss

- **Issue:** Improper upgrade procedures can corrupt or delete data.
- **Solution:** Always back up using pg_dump, pg_dumpall, or filesystem snapshots. Test restore procedures.

5. Performance Regressions

- **Issue:** New versions may introduce changes that degrade performance for certain workloads.
- **Solution:** Benchmark queries before and after upgrade. Tune configuration parameters for the new version.

Benefits of Upgrading PostgreSQL

Upgrading unlocks powerful features and performance improvements. For example, PostgreSQL 13 introduced:



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

- **Improved Query Planning:** Better statistics for complex filters.
- **Parallel Vacuuming:** Faster table maintenance.

VACUUM (PARALLEL 4) customers;

- **Incremental Sorting:** Efficient multi-key sorting.

SELECT name FROM employees ORDER BY salary DESC, name ASC;

Step-by-Step Upgrade Guide

Here's a high-level upgrade workflow:

◆ Preparation

1. Read release notes
2. Audit extensions and custom functions
3. Back up your data:

`pg_dumpall > alldb.bak`

◆ Installation

4. Install new PostgreSQL binaries
5. Initialize new cluster:

`initdb -D /new/data`

◆ Migration

6. Use pg_upgrade:

`pg_upgrade -b /old/bin -B /new/bin -d /old/data -D /new/data`

7. Run post-upgrade scripts:

`./analyze_new_cluster.sh`

`./delete_old_cluster.sh`

◆ Validation

8. Test application behavior
9. Monitor logs and performance
10. Keep old version temporarily for rollback

Recommended Guides



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

- Percona: Common PostgreSQL Upgrade Errors and How to Avoid Them
- Reintech: Best Practices and Pitfalls
- SQL-Easy: Comprehensive Upgrade Guide

I used ClusterControl to deploy a PostgreSQL 11 database cluster for a single node. The following have been tested on Centos 7 and Ubuntu Focal (20.04.1):

```
$ /usr/pgsql-11/bin/postgres --version
```

```
postgres (PostgreSQL) 11.13
```

```
postgres=# dx
```

List of installed extensions

Name	Version	Schema	Description
fuzzystrmatch	1.1	public	determine similarities and distance between strings
pg_stat_statements	1.6	public	track execution statistics of all SQL statements executed
plpgsql	1.0	pg_catalog	PL/pgSQL procedural language
postgis	3.1.4	public	PostGIS geometry and geography spatial types and functions
postgis_raster	3.1.4	public	PostGIS raster types and functions
postgis_sfsgal	3.1.4	public	PostGIS SFSGAL functions
postgis_tiger_geocoder	3.1.4	tiger	PostGIS tiger geocoder and reverse geocoder
postgis_topology	3.1.4	topology	PostGIS topology spatial types and functions
timescaledb	2.3.1	public	Enables scalable inserts and complex queries for time-series data

(9 rows)



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

So I got the following,

PostgreSQL server version: 11.13

TimescaleDB version: 2.3.1

PostGIS version: 3.1.4

If you want to test this with ClusterControl, there are two ways to have TimescaleDB.

Setup the required repositories

First, let's add the PostgreSQL repository.

For CentOS/RHEL/Oracle Linux

You have to make sure that you have the right repository. For Enterprise Linux (EL) 7, you can do the following:

```
sudo yum -y install https://download.postgresql.org/pub/repos/yum/reporpms/EL-7-x86_64/pgdg-redhat-repo-latest.noarch.rpm
```

Let's add the TimescaleDB repository as well.

```
vi /etc/yum.repos.d/timescale_timescaledb.repo
```

Then add the following contents for this file,

```
[timescale_timescaledb]
```

```
name=timescale_timescaledb
```

```
baseurl=https://packagecloud.io/timescale/timescaledb/el/7/$basearch
```

```
repo_gpgcheck=1
```



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

gpgcheck=0

enabled=1

gpgkey=https://packagecloud.io/timescale/timescaledb/gpgkey

sslverify=1

sslcacert=/etc/pki/tls/certs/ca-bundle.crt

metadata_expire=300

Just replace the baseurl accordingly if you are using a version other than EL 7.

For Ubuntu/Debian

Add the PG repository for Ubuntu Focal:

```
deb http://apt.postgresql.org/pub/repos/apt/ focal-pgdg main
```

For other Ubuntu/Debian distributions, just replace the focal accordingly, which you can find here <http://apt.postgresql.org/pub/repos/apt/dists/>. For example, replace focal-pgdg with buster-pgdg.

Now, let's add the repository for TimescaleDB,

```
sudo sh -c "echo 'deb [signed-by=/usr/share/keyrings/timescale.keyring]
https://packagecloud.io/timescale/timescaledb/ubuntu/ $(lsb_release -c -s) main' >
/etc/apt/sources.list.d/timescaledb.list"
```

Import the keyring,

```
wget --quiet -O - https://packagecloud.io/timescale/timescaledb/gpgkey | sudo gpg --dearmor -o /usr/share/keyrings/timescale.keyring
```

and update the package lists for upgrades for packages that need upgrading, as well as new packages that have just come to the repositories.

```
sudo apt-get update
```

You can replace the sub-directory in the URL if you are using Debian from ubuntu.

Now that we have the repository ready, we're good to go.

Install PostgreSQL version 13 with TimescaleDB and PostGIS



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

Installing PostgreSQL 13 can be done on the same host. First, you must make sure things such as the database port are unique. In other words, it has to be different from the current PostgreSQL 11 installed on the same host.

For CentOS/RHEL/Oracle Linux

Run the command below to install PostgreSQL 13 and its dependent packages:

```
yum install postgresql13.x86_64 postgresql13-server.x86_64 postgresql13-contrib.x86_64 postgresql13-libs.x86_64
```

Then initialize the database cluster and its required collection of databases by running the command below:

```
$ /usr/pgsql-13/bin/postgresql-13-setup initdb
```

At this point, there should be two data directories both for PG 11 and PG 13:

```
[root@pupnode28 ~]# ls -alth /var/lib/pgsql/* -d
drwx----- 4 postgres postgres 51 Sep 22 14:19 /var/lib/pgsql/13
drwx----- 4 postgres postgres 33 Sep 21 18:53 /var/lib/pgsql/11
```

Now that we're good with PostgreSQL 13 let's install TimescaleDB. We need to make sure that the plugin to be installed is the same version on PostgreSQL 11.

Take note that, in order to make sure that pg_upgrade will work smoothly, the plugins of your source and major version destination should be the same version. This is because pg_upgrade will look for its designated libraries linked to the plugins or extensions that have been loaded or used by your old or source database version of your PostgreSQL. You can verify this in your Enterprise Linux by running showduplicates or by verifying with info just like below either with dnf or yum:

```
$ yum --showduplicates list timescaledb_13.x86_64 timescaledb-2-postgresql-13.x86_64 timescaledb-2-oss-postgresql-13.x86_64 timescaledb-2-loader-postgresql-13.x86_64|grep '2.3.1'
```

Repository pgdg-common is listed more than once in the configuration

```
timescaledb-2-loader-postgresql-13.x86_64 2.3.1-0.el7    timescale_timescaledb
timescaledb-2-oss-postgresql-13.x86_64  2.3.1-0.el7    timescale_timescaledb
timescaledb-2-postgresql-13.x86_64     2.3.1-0.el7    timescale_timescaledb
```

Or verify it with the info option:



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

```
$ yum info timescaledb-2-loader-postgresql-13-2.3.1-0.el7.x86_64 timescaledb-2-postgresql-13-2.3.1-0.el7.x86_64
```

Now, we're ready to install the TimescaleDB package for the PG 13 version.

```
$ yum install timescaledb-2-loader-postgresql-13-2.3.1-0.el7.x86_64 timescaledb-2-postgresql-13-2.3.1-0.el7.x86_64
```

After you've installed it, you can try to run timescaledb-tune tool to tune up your postgresql.conf configuration file. Just run the command below:

```
$ timescaledb-tune --pg-config=/usr/pgsql-13/bin/pg_config
```

Now, let's install the PostGIS package for the PG 13 version as well.

```
$ yum install -y postgis31_13.x86_64
```

For Ubuntu/Debian

Simply run:

```
$ apt install postgresql-client-13 postgresql-13
```

The great thing with Ubuntu/Debian distributions is that there are tools for PostgreSQL that are very handy for managing your PostgreSQL clusters, such as pg_lsclusters, pg_ctlcluster, etc.

You can verify your available clusters installed.

```
$ pg_lsclusters
```

Ver	Cluster	Port	Status	Owner	Data directory	Log file
-----	---------	------	--------	-------	----------------	----------

11	main	5432	online	postgres	/var/lib/postgresql/11/main	log/postgresql-%Y-%m-%d_%H%M%S.log
----	------	------	--------	----------	-----------------------------	------------------------------------

13	main	5433	online	postgres	/var/lib/postgresql/13/main	
						/var/log/postgresql/postgresql-13-main.log

In Ubuntu/Debian, there's no need to change the port since it will be handled during the installation phase and it detects and sets it uniquely, accordingly.

Now, let's install TimescaleDB.

```
$ apt install timescaledb-2-loader-postgresql-13 timescaledb-2-postgresql-13
```

Optionally, you can run the timescaledb-tune tool to tune up your postgresql.conf configuration file by simply just invoking the tool as the following:



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

\$ timescaledb-tune

Now, we're ready to install the PostGIS package for PG 13.

```
$ apt install postgresql-13-postgis-3-scripts postgresql-13-postgis-3
```

Review your postgresql.conf

It's always better to review your postgresql.conf configuration file. In Enterprise Linux versions, you can locate your postgresql.conf either in your data_directory or PGDATA path. Whereas for Ubuntu/Debian, you can locate it in /etc/postgresql///postgresql.conf. Make sure that in your postgresql.conf, the following lines are correctly configured:

```
shared_preload_libraries = 'pg_stat_statements,timescaledb' # pg_stat_statements is not required but if you are using ClusterControl, make sure this is appended.
```

```
port = 5532 # make sure that the port number is unique than the old version of your PostgreSQL
```

```
listen_address = * # depends on your setup but if you need to specify the available network interfaces to its IP addresses (IPv4 or IPv6) set it accordingly.
```

It's best practice to compare your old and new versions of your PostgreSQL configuration files to ensure that your postgresql.conf is identical to what is needed and set.

Before proceeding with the next step, we also need to check that your PostgreSQL version 13 is loaded accordingly. Make sure that you have the latest version acquired or installed in your host. Start the database and make sure that it starts and runs correctly.

To start in EL distributions, run the command below:

```
$ sudo -iu postgres /usr/pgsql-13/bin/pg_ctl -D /var/lib/pgsql/13/main/ -o "-c config_file=/var/lib/pgsql/13/main/postgresql.conf" start
```

or for Ubuntu/Debian, run the command below:

```
$ sudo -iu postgres /usr/lib/postgresql/13/bin/pg_ctl -D /var/lib/postgresql/13/main/ -o "-c config_file=/etc/postgresql/13/main/postgresql.conf" start
```

or use the pg_ctlcluster tool to start, restart, or stop your PG Cluster.



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

When ready, run pg_upgrade...

Before moving on, first make sure that you always have your backup from your old server ready and available. Always take a logical backup and physical backup as good practice before proceeding with a major upgrade.

Now that you're ready, you are good to run pg_upgrade. In practice, you have to initially run pg_upgrade with a check to determine the incompatibility and issues before proceeding to the main procedure of pg_upgrade. Before running the pg_upgrade, make sure that both PG 11 and PG 13 are down while doing this process. That just means downtime is needed for this process.

To do that, run the following command with –check option:

```
$ sudo -iu postgres /usr/lib/postgresql/13/bin/pg_upgrade -o "-c config_file=/etc/postgresql/11/main/postgresql.conf" --old-datadir=/var/lib/postgresql/11/main/ -O "-c config_file=/etc/postgresql/13/main/postgresql.conf" --new-datadir=/var/lib/postgresql/13/main/ --old-bindir=/usr/lib/postgresql/11/bin --new-bindir=/usr/lib/postgresql/13/bin --check
```

Replace the –old-datadir value or config_file accordingly if it's different from your setup.

This shall run compatibility checks just like the result below:

Performing Consistency Checks

```
-----  
Checking cluster versions          ok  
Checking database user is the install user      ok  
Checking database connection settings      ok  
Checking for prepared transactions      ok  
Checking for system-defined composite types in user tables  ok  
Checking for reg* data types in user tables      ok  
Checking for contrib/isn with bigint-passing mismatch    ok  
Checking for tables WITH OIDS          ok  
Checking for invalid "sql_identifier" user columns      ok  
Checking for presence of required libraries      ok  
Checking database user is the install user      ok
```



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

Checking for prepared transactions	ok
Checking for new cluster tablespace directories	ok

Clusters are compatible

If all checks signal “ok,” that means a successful check, and the bottom message shows that clusters are compatible, then you should be good to go.

Finally, run the command again without the –check option:

```
$ sudo -iu postgres /usr/lib/postgresql/13/bin/pg_upgrade -o "-c config_file=/etc/postgresql/11/main/postgresql.conf" --old-datadir=/var/lib/postgresql/11/main/ -O "-c config_file=/etc/postgresql/13/main/postgresql.conf" --new-datadir=/var/lib/postgresql/13/main/ --old-bindir=/usr/lib/postgresql/11/bin --new-bindir=/usr/lib/postgresql/13/bin
```

Performing Consistency Checks

Checking cluster versions	ok
Checking database user is the install user	ok
Checking database connection settings	ok
Checking for prepared transactions	ok
Checking for system-defined composite types in user tables	ok
Checking for reg* data types in user tables	ok
Checking for contrib/isn with bigint-passing mismatch	ok
Checking for tables WITH OIDS	ok
Checking for invalid "sql_identifier" user columns	ok
Creating dump of global objects	ok
Creating dump of database schemas	ok
Checking for presence of required libraries	ok



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

Checking database user is the install user ok

Checking for prepared transactions ok

Checking for new cluster tablespace directories ok

If pg_upgrade fails after this point, you must re-initdb the

new cluster before continuing.

Performing Upgrade

Analyzing all rows in the new cluster ok

Freezing all rows in the new cluster ok

Deleting files from new pg_xact ok

Copying old pg_xact to new server ok

Setting oldest XID for new cluster ok

Setting next transaction ID and epoch for new cluster ok

Deleting files from new pg_multixact/offsets ok

Copying old pg_multixact/offsets to new server ok

Deleting files from new pg_multixact/members ok

Copying old pg_multixact/members to new server ok

Setting next multixact ID and offset for new cluster ok

Resetting WAL archives ok

Setting frozenxid and minmxid counters in new cluster ok

Restoring global objects in the new cluster ok

Restoring database schemas in the new cluster ok

Copying user relation files ok

Setting next OID for new cluster ok

Sync data directory to disk ok

Creating script to analyze new cluster ok



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

Creating script to delete old cluster ok

Checking for extension updates notice

Your installation contains extensions that should be updated
with the ALTER EXTENSION command. The file

`update_extensions.sql`

when executed by psql by the database superuser will update
these extensions.

Upgrade Complete

Optimizer statistics are not transferred by pg_upgrade so,
once you start the new server, consider running:

`./analyze_new_cluster.sh`

Running this script will delete the old cluster's data files:

`./delete_old_cluster.sh`

Depending on the size of your dataset, it might take a bit of time. In the example command above, it does copy the transaction logs, which are of physical files. However, if your disk size is tight, then you can use the -k or –link option, which uses hard links. If all goes well, it should provide you with messages such as the above notifying you with complete upgrade and notices to update the extensions you might have. In this setup, it goes smoothly. The `update_extensions.sql` contains only the following:

```
$ cat update_extensions.sql
connect postgres
ALTER EXTENSION "pg_stat_statements" UPDATE;
```

As you noticed, pg_upgrade does a series of actions. It analyzes all rows from the source cluster, copying transaction metadata logs and its multi-transaction status data and the rest.

Once you figured out that everything looks good, then run the `analyze_new_cluster.sh` script, which shall run



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

vacuumdb --all --analyze-only.

```
$ sudo -iu postgres PGPORT=5433 ./analyze_new_cluster.sh
```

Take note that you should specify the port of your PostgreSQL 13 so that vacuumdb will run correctly to the right target server.

In this case, the upgrade outcome with TimescaleDB and PostGIS extensions enabled goes well. Once all appears in order, make sure to start the server and do a series of tests and checks.

Test the upgrade process

Always test and review the upgrade process. Here are a few tables and a user-defined database, which contains hypertables and PostGIS tables that rely on using geometry and geography spatial types and functions.

```
$ sudo -iu postgres psql -p5433
```

```
psql (13.4 (Ubuntu 13.4-1.pgdg20.04+1))
```

Type "help" for help.

```
postgres=# \l
```

List of databases

Name	Owner	Encoding	Collate	Ctype	Access privileges
mydb	postgres	UTF8	C.UTF-8	C.UTF-8	
postgres	postgres	UTF8	C.UTF-8	C.UTF-8	
template0	postgres	UTF8	C.UTF-8	C.UTF-8	=c/postgres + postgres=CTc/postgres
template1	postgres	UTF8	C.UTF-8	C.UTF-8	postgres=CTc/postgres+ =c/postgres

(4 rows)

```
postgres=# c mydb
```

You are now connected to database "mydb" as user "postgres".

```
mydb=# set search_path="$user",public;
```



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

SET

mydb=# dt+

List of relations

Schema	Name	Type	Owner	Persistence	Size	Description
--------	------	------	-------	-------------	------	-------------

public	conditions	table	postgres	permanent	8192 bytes	
--------	------------	-------	----------	-----------	------------	--

public	global_points	table	postgres	permanent	16 kB	
--------	---------------	-------	----------	-----------	-------	--

public	roads	table	postgres	permanent	16 kB	
--------	-------	-------	----------	-----------	-------	--

public	sample_table	table	postgres	permanent	8192 bytes	
--------	--------------	-------	----------	-----------	------------	--

public	spatial_ref_sys	table	postgres	permanent	6968 kB	
--------	-----------------	-------	----------	-----------	---------	--

(5 rows)

Checking some of my PostGIS and hypertables:

mydb=# d roads

Table "public.roads"

Column	Type	Collation	Nullable	Default
--------	------	-----------	----------	---------

road_id	integer			
---------	---------	--	--	--

road_name	character varying			
-----------	-------------------	--	--	--

roads_geom	geometry(LineString)			
------------	----------------------	--	--	--

mydb=# d sample_table

Table "public.sample_table"

Column	Type	Collation	Nullable	Default
--------	------	-----------	----------	---------

id	integer		not null	nextval('sample_table_id_seq'::regclass)
----	---------	--	----------	--



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

time | timestamp with time zone | | not null |

name | character varying | | not null |

Indexes:

"sample_table_pkey" PRIMARY KEY, btree (id, "time")

"sample_table_time_idx" btree ("time" DESC)

Triggers:

ts_insert_blocker BEFORE INSERT ON sample_table FOR EACH ROW EXECUTE
FUNCTION _timescaledb_internal.insert_blocker()

Number of child tables: 371 (Use d+ to list them.)

mydb=# d conditions

Table "public.conditions"

Column	Type	Collation	Nullable	Default
--------	------	-----------	----------	---------

time	timestamp with time zone		not null	
------	--------------------------	--	----------	--

location	text		not null	
----------	------	--	----------	--

location2	character(10)		not null	
-----------	---------------	--	----------	--

temperature	double precision			
-------------	------------------	--	--	--

humidity	double precision			
----------	------------------	--	--	--

Indexes:

"conditions_time_idx" btree ("time" DESC)

Triggers:

ts_insert_blocker BEFORE INSERT ON conditions FOR EACH ROW EXECUTE
FUNCTION _timescaledb_internal.insert_blocker()

Number of child tables: 366 (Use d+ to list them.)

mydb=# select count(*) from sample_table;



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

count

2588

(1 row)

```
mydb=# SELECT name FROM global_points WHERE ST_DWithin(location,  
'SRID=4326;POINT(-110 29)::geography, 1000000);
```

name

Town

Forest

(2 rows)

```
mydb=# SELECT n, ST_AsEWKT(ST_GeometryN(the_geom, n)) As geomewkt
```

```
mydb# FROM (
```

```
mydb# VALUES (ST_GeomFromEWKT('MULTIPOINT(1 2 7, 3 4 7, 5 6 7, 8 9 10)'),
```

```
mydb# ( ST_GeomFromEWKT('MULTICURVE(CIRCULARSTRING(2.5 2.5,4.5 2.5, 3.5  
3.5), (10 11, 12 11))'))
```

```
mydb# )As foo(the_geom)
```

```
mydb# CROSS JOIN generate_series(1,100) n
```

```
mydb# WHERE n <= ST_NumGeometries(the_geom);
```

```
n |      geomewkt
```

```
--+-----
```

```
1 | POINT(1 2 7)
```

```
1| CIRCULARSTRING(2.5 2.5,4.5 2.5,3.5 3.5)
```

```
2 | POINT(3 4 7)
```

```
2 | LINESTRING(10 11,12 11)
```

```
3 | POINT(5 6 7)
```

```
4 | POINT(8 9 10)
```

(6 rows)



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

Now everything looks ready to serve as my new cluster.

Once you get things going, then you can run:

```
$ sudo -iu postgres ./delete_old_cluster.sh
```

Make sure you only run the script since this is a very dangerous script, as it will delete all files in your old PostgreSQL cluster.

Cluster overview


The setup involves a **2-node PostgreSQL cluster** using [Zalando's Spilo](#), which bundles:

- **PostgreSQL + TimescaleDB** — For time-series data storage.
- [**Patroni**](#) — Manages the cluster state, handles failover, and ensures high availability.

Each node is either a **Leader** or a **Replica**, and both communicate with a set of **ETCD nodes** that serve as a distributed key-value store for cluster coordination.

Here's a high-level architecture diagram to visualize the setup:

Press enter or click to view image in full size



High-level overview of PostgreSQL cluster using Spilo

Note: While this example uses only 2 nodes for simplicity, the setup is extendable to larger clusters.

With that in place, let's explore the upgrade process step by step.

Checklist — If You're Doing This in a Production Setup

If you're planning a **major version upgrade** in a live environment, several precautions are critical — such as validating functionality, choosing the right upgrade window, and informing stakeholders.

Rather than repeating the full **checklist** here, I recommend referring to the [Minor Version Upgrade Guide](#), which outlines a comprehensive pre-upgrade checklist suitable for both minor and major upgrades.

The only major difference: expect **longer downtime and more complexity** with a major version upgrade — especially for replica nodes.

High-level overview

This section provides a high-level overview of the process without going into technical details. Detailed technical explanations are provided later in the post.

Major version upgrades involve distinct steps for leader and replica nodes, including:

1. Pre-upgrade processing.
2. Upgrade a leader node.
3. Upgrade a replica node.



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

4. Post-upgrade processing.

Pre-upgrade processing

Before upgrading, ensure that replica nodes can communicate with the leader node via SSH, and have an additional node for complete rollback or failure recovery.

Upgrade a leader node

Utilize PostgreSQL's [pg_upgrade](#) tool for the leader node upgrade. The process involves the following steps:

1. Disable all write activity on the leader node.
2. Temporarily [step down](#) patroni from managing the cluster.
3. Shutdown database (or Postgres service).
4. Initialize DB ([initdb](#)).
5. Copy configurations such as pg_hba.conf, postgres.conf etc.
6. Perform upgrade ([pg_upgrade](#)).
7. Please check the output of the upgrade for any warnings, it might guide you on things you need to take care of after the upgrade.
8. Adjust patroni configuration.
9. Wipe complete cluster state from DCS (Dynamic Configuration Settings).
10. Start Postgres service with target version.
11. Start patroni.

Congrats 🎉, you have successfully upgraded the leader node.

Upgrade a replica node

Unlike the leader node, replicas cannot use [pg_upgrade](#). The basic idea for upgrading an existing replica node is to synchronize folder structure (data & Write Ahead Log (WAL) directories).

1. Shut down database (or Postgres service).
2. Synchronize data and Write Ahead Log (WAL) directories.
3. Adjust patroni configuration
4. Start up database (or Postgres service).
5. Configure replication.



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

Congrats 🎉, you have successfully upgraded the replica node.

Post-upgrade processing

Enable write activity on the leader node and rebuild statistics for tables.

Technical details

Close to zero downtime using in-place upgrade

Major version upgrades require downtime, but in-place upgrades minimize this downtime. Using pg_upgrade with the --link option reduces downtime to near zero by copying only metadata.


Tools	**Reloading data**	**Downtime needed**
`pg_upgrade` Metadata and files are copied	Downtime is needed	
`pg_upgrade --link` Only metadata are copied	Close to zero	

Source: [Upgrading and updating PostgreSQL](#)

Replica's Upgrade Process — Synchronization of Folder Structure

You can understand this concept if you are familiar with replication.

Press enter or click to view image in full size



High-level overview of a replication



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

When incorporating a new replica node into our cluster, the leader node performs the following steps:

1. **Initial Data Replication:** The leader node replicates everything except for WAL segments to the new replica.
2. **WAL Segment Replication:** Subsequently, it replicates WAL segments to the replica.

Synchronizing folder structures manually performs replication, effectively upgrading existing replica nodes.

Ensure Replica-Node Communication with Leader via SSH

Use rsync for folder synchronization, requiring SSH communication between replica and leader nodes.

You can use the following script to ensure successful communication.

```
#!/bin/bash

# Configuration
LEADER_IP="..."
LEADER_USER="..."

REPLICA_IP="..."
REPLICA_USER="..."

# generate key
ssh -i ~/.ssh/your.pem "$REPLICA_USER@$REPLICA_IP" \
'sudo ssh-keygen -t rsa -b 2048 -f /root/.ssh/id_rsa -N ""

# download generated key on your machine
rsync -avz --progress --rsync-path="sudo rsync" \
-e "ssh -i ~/.ssh/your.pem" \
"$REPLICA_USER@$REPLICA_IP:/root/.ssh/id_rsa.pub" "$PWD/replica-ssh-key.pub"

# copying it to a leader node
cat replica-ssh-key.pub | ssh -i ~/.ssh/your.pem "$LEADER_USER@$LEADER_IP" \
'cat >> /home/ec2-user/.ssh/authorized_keys && echo "Key copied"'
```

⚠ You need to run this script from your machine.

Importance of Temporarily Stepping Down Patroni



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

Patroni manages cluster. It keeps track of a leader's health and act accordingly. It will perform a failover when a leader is unhealthy to ensure high availability.

As upgrading process makes leader unhealthy, we let Patroni to step down temporarily and don't do anything, and we're aware of the situation.

Upgrade in Action

This section provides practical insights and applies the aforementioned workflow to enhance understanding.

Source code: <https://github.com/anasanjaria/blogs/tree/main/postgres-upgrade>

For demonstration, assume a cluster of two nodes: one leader and one replica, running PostgreSQL version 13.x and upgrading to version 15.x.

For the sake of simlicity, I am running all the containers locally and persisting data on different directories.

Data for **node-1** & **node-2** are stored in **data-1** & **data-2** respectively.

Upgrading leader node

Execute the provided script to upgrade the leader node. The script guides through the process, requiring user input at various steps.

```
docker exec -it node-1 /opt/pgupgrade.sh
```

Upgrading replica node

Similarly, use the provided script to upgrade an existing replica node. Execute the script and follow the prompts for a successful upgrade.

```
./pgupgrade-replica.sh
```

Post upgrade process

Enabling write activity — Let's deactivate write lock now.

```
curl --request PATCH \
--url http://localhost:8008/config \
--header 'Content-Type: application/json' \
--data '{
"postgresql": {
"parameters": {
"default_transaction_read_only": "off"
}
}
```



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

}

'

Rebuilding statistics — Rebuilding statistics is extremely crucial to have better performance of your production system.

If you skip this step, you will end up with extremely slow queries.

Typically, analyze command rebuilds statistics, but it's extremely slow. Hence, it's not practical for big databases.

vacuumdb is another tool for the same purpose and also supports parallelism. Hence, the preferred way for production system.

```
/path/to/vacuumdb -d <DATABASE> \
-t <TABLE-1> -t <TABLE-2> ... -t <TABLE-n> \
--analyze-in-stages \
-j <CONCURRENT_CONNECTIONS>
```

TIP: Use the same number of connections as the number of tables because a single connection is used per table.

Conclusion

The successful upgrade of our production cluster marks a significant milestone in our journey towards maintaining a robust and reliable infrastructure.

By extending the upgrade process to include both leader and replica nodes, we've not only mitigated potential risks but also enhanced the stability and resilience of our critical services.

I hope you've found this journey as enlightening and empowering as I have. Here's to smooth upgrades, robust clusters, and endless possibilities!



Mukesh Chaurasia – SQL DBA


COURSE: POSTGRESQL DATABASE ADMINISTRATOR

How to Uninstall PostgreSQL

Steps


Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR



1.

Prepare. Before you start to uninstall "P - SQL", you may need to save your personal settings and data files: open up pgAdmin III - Object browser, visit the Tools menu; scroll down and run the option "Backup...".




2.

Launch the option "Chang/Remove" for PostgreSQL program, version 9.1 (size, 172.00MB).


Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR



3.

Click on the "Yes" button on the "Question" window and then wait.




Click OK on Warning. You may then manually delete leftovers related to PostgreSQL, in its installation directory.



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR



5.

You then click "OK" on "Info".




6.

Make sure that you have saved all of your files and programs. You then click Yes in Question information box to restart your device.


Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR



7.

Remove the pgAgent 3.0.1 (size, 9.48MB).




8.

Again, you may clean the file leftovers by hand.


Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR



9.

Continue to remove MS Visual C++ 2008, a component for PostgreSQL software.



10.

Remove C++ 2005 project listed.



11.

Save and reboot your computer. You manually search and clean those traces related the programs installed above.

How to Uninstall PostgreSQL From Windows

What is the Purpose to Uninstall Postgres on Windows?

Uninstalling Postgres ensures that all of its files, configurations, and services are removed from the computer, freeing up disk space and resources. Additionally, it is useful if users are experiencing issues with Postgres and need to perform a clean installation to troubleshoot the problem.

Method: Using Control Panel to Uninstall Postgres


Uninstalling PostgreSQL on Windows can be done through the Control Panel or by using the PostgreSQL installer. Here are the steps to follow:

Step 1: Open the Control Panel

Go to the Windows Start menu and search for Control Panel, then click on it to open it:


Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR



Step 2: Go to “Programs and Features”


Once you reach the Control Panel, click on “**Programs and Features**” (or "Uninstall a program" depending on the version of Windows):



Step 3: Locate PostgreSQL


Scroll down the list of programs until you find PostgreSQL. Click on it to select it:

COURSE: POSTGRESQL DATABASE ADMINISTRATOR



Step 4: Uninstall PostgreSQL


Click on the “Uninstall” button by pressing the right button of the mouse on the “PostgreSQL” program:




Follow the prompts to complete the uninstallation process:

Mukesh Chaurasia – SQL DBA


COURSE: POSTGRESQL DATABASE ADMINISTRATOR



Hit the “Next” button to begin the uninstallation of Postgres:




After completing the uninstallation, a new window will appear:



Hit the “Ok” button to finish.

Note: Once the uninstallation is done, the user can manually delete the PostgreSQL files and folders that were left behind. These can typically be found in the “**Program Files**” or “**Program Files (x86)**” folder in the Windows installation:



Remove the “PostgreSQL” directory to completely uninstall Postgres from your Windows system.



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

How To Remove PostgreSQL in UNIX

Step 1: List the PostgreSQL Packages

Use the dpkg tool to list packages pertaining to the PostgreSQL setup.

`dpkg -l | grep postgres`

Output:

```
root@newclient:~# dpkg -l | grep postgres
ii  pgdg-keyring                         2018.2          all      keyring for apt.postgresql.org
ii  postgresql-10                          10.10-1.pgdg18.04+1   amd64   object-relational SQL database, versio
n 10 server
ii  postgresql-client-10                   10.10-1.pgdg18.04+1   amd64   front-end programs for PostgreSQL 10
ii  postgresql-client-common               204.pgdg18.04+1     all     manager for multiple PostgreSQL client
versions
ii  postgresql-common                     204.pgdg18.04+1     all     PostgreSQL database-cluster manager
root@newclient:~#
```

Step 2: Delete the PostgreSQL Packages

In Step 1, all of the software packages related to the PostgreSQL install are shown. To remove and delete them all with one command, Start with this command.

`apt-get --purge remove` command

This is followed by each package name indicated separated with a space. For my particular version of PostgreSQL, the software installed was:

- pgdg-keyring
- postgresql-10
- postgresql-client-10
- postgresql-client-common
- postgresql-common

so, my purge remove command will look like this:

```
sudo apt-get --purge remove pgdg-keyring postgresql-10 postgresql-client-10
postgresql-client-common postgresql-common
```

Step 3: Verifying the Deletion of PostgreSQL

Once you remove these packages you should no longer be able to enter into the PostgreSQL environment. You can verify that by running the grep command again and searching for postgres:

```
root@newclient:~# dpkg -l | grep postgres
```

```
root@newclient:~#
```



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

As you can see, there is no output from that command which means, PostgreSQL has been successfully uninstalled!

Removing and Reinstalling PostgreSQL

Step-by-Step Guide to Remove PostgreSQL

Stop PostgreSQL Services

Before uninstalling PostgreSQL, stop all running PostgreSQL processes. To check for active PostgreSQL processes, use this command:

```
ps aux | grep postgres
```

If any PostgreSQL processes are running, stop them with:

```
sudo systemctl stop postgresql
```

Tip: Verify PostgreSQL Service Status

After stopping the PostgreSQL service, you can verify its status using:

```
sudo systemctl status postgresql
```

This command will show you if the service is inactive or if there are any remaining processes.

Uninstall PostgreSQL Packages

To remove PostgreSQL packages, use this command:

```
sudo apt-get remove --purge postgresql*
```

This command removes all PostgreSQL packages and their config files.

Clean Up Remaining Files and Directories

After uninstalling the packages, you may need to remove some leftover directories:

```
sudo rm -rf /etc/postgresql/
```

```
sudo rm -rf /var/lib/postgresql/
```

```
sudo rm -rf /var/log/postgresql/
```

These commands delete the PostgreSQL config, data, and log directories.

Remove PostgreSQL User and Group

Remove the PostgreSQL system user and group:



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

`sudo userdel postgres`

`sudo groupdel postgres`

These commands delete the PostgreSQL user and group from the system.

Reinstalling PostgreSQL on Ubuntu

Preparing for a Fresh Installation

Before reinstalling PostgreSQL, update your package lists:

`sudo apt update`

Check that your system meets the requirements for PostgreSQL installation. Ubuntu usually includes all needed dependencies.

Tip: Check PostgreSQL Version

To check the available PostgreSQL versions in your repositories, use:

`apt-cache search postgresql | grep postgresql`

This command lists all available PostgreSQL packages and their versions.

Installing PostgreSQL

To install PostgreSQL, run:

`sudo apt install postgresql`

This installs the latest version of PostgreSQL from the Ubuntu repositories.

To install a specific PostgreSQL version, use:

`sudo apt install postgresql-<version>`

Replace <version> with the PostgreSQL version number you want (e.g., 12, 13, 14).

Configuring the New PostgreSQL Installation

After installation, PostgreSQL creates a default database and user. To set up more databases, use:

`sudo -u postgres createdb <database_name>`

To create a new user role:

`sudo -u postgres createuser --interactive`

This starts a prompt to set up a new user role.

To set user permissions, access the PostgreSQL prompt:



Mukesh Chaurasia – SQL DBA

COURSE: POSTGRESQL DATABASE ADMINISTRATOR

sudo -u postgres psql

Then, use SQL commands to give privileges:

GRANT ALL PRIVILEGES ON DATABASE <database_name> TO <username>;

Replace <database_name> and <username> with your database and user names.
